

# Data Tidying and Merging

Juan Carlos Villaseñor-Derbez (JC)

Last week you were asked:

How much money have tuna purse seiners made since 2000 when fishing for bigeye tuna (*Thunnus obesus*) in the Eastern Pacific Ocean?

We made some simplifying assumptions and got some values (a total of 3,070 M USD since 2000, or about 127 M USD per year). You are now tasked with coming up with more refined estimates. For example, we will account for the fact that the price of fish varies every year.

How we will approach this:

- Find data that shows prices per year and species
- Read them, clean them, tidy them up (The “data tidying” part)
- Combine our catch data from last week with this new price data (The “merging” part)
- Re-calculate our total revenues since 2000

This will require three pipelines: - Tidy price data (Exercise 1) - Wrangle catch data (Exercise 2) - Combine tidy prices and catch data (Exercise 3)

Pipelines 1 and 3 contain tools covered this week. You should already be familiar with pipeline 2.

## Exercise 1: Tidying price data

### Part A: Downloading the data

#### Post-it up

1. In a web browser, go to [ffa.int](http://ffa.int). This is the website for the Pacific Islands Forum Fisheries Agency
2. Hover over “Publication and Statistics” on the top menu
3. Select “Statistics”
4. You will be taken to a site with five items. Download the zip folder called **Economic and Development Indicators and Statistics: Tuna Fishery of the Western and Central Pacific Ocean 2024**
5. As before, place the downloaded zip file in your EVR628/data/raw folder and proceed to extract it

6. Open the excel file called **Compendium of Economic and Development Statistics 2024** and study the **Contents** tab
7. Can you identify the price data that we need?
  - Which sheet
  - What range?

**Post-it down**

## Part B: Reading excel data

**Post-it up**

1. Open your RStudio project for EVR628
2. In your console, install the `readxl` package: `install.packages("readxl")`
3. Start a **new** script called `tuna_analysis_prices.R`<sup>1</sup>
4. Add the usual code commenting outline
5. We will need three packages: `readxl`, `janitor`, and `tidyverse`, load them at the top of your script using `library()`
6. Use `?read_excel()` to look at the documentation for the function
7. Use `read_excel()` to create a new object called `tuna_prices` and read the price data we need<sup>2</sup>. Immediately pipe it into `clean_names`.

**Post-it down**

## Part C: Inspecting price data

Be prepared to discuss the following points:

**Post-it up**

1. Inspect the column names of `tuna_prices` using `colnames()` in your console.
2. How many columns and rows do we have?
3. Any missing values?
4. Do we need to make the data wider or longer?
5. Using comments, write out what the target data should be (expand my code chunk see what I wrote)

**Post-it down**

See an example of my description below

```
[1] "year"           "japan_fresha"   "japan_frozenb" "us_freshc"
[5] "us_frozend"
[1] 28  5
```

---

<sup>1</sup>I would typically suggest to overwrite whatever we had last week in `tuna_analysis.R` because GitHub would keep a version, but I understand you might want to keep the script as is

<sup>2</sup>Hint: You will need to specify a file path, a sheet, and a range of cells.

```

# A tibble: 4 x 5
  year japan_fresha japan_frozenb us_freshc us_frozend
  <dbl>      <dbl>      <dbl>      <dbl>      <dbl>
1 1997       8204.     8169.      NA        NA
2 1998       7703.     6320.      NA        NA
3 1999       8809.     9093.      NA        NA
4 2000       9198.     8557.      NA        NA
# The final data set should have two columns: year and price. Since we have four
# prices (two markets, two presentations), I will use the average price per year.
# The tidy data set should therefore have four columns: year, market,
# presentation, and price.

```

## Part D: Tidy your price data

### Post-it up

1. Look at the documentation for your `pivot_*` function. What does it say about cases where `names_to` is of length  $> 1$ ?
2. What about the `names_sep` argument?
3. Use the appropriate `pivot_*` function to reshape your data and save them to a new object called `tidy_tuna_prices`<sup>3</sup>
4. Your resulting tibble should have 104 rows and 4 columns and look like this:<sup>4</sup>

```

# A tibble: 104 x 4
  year market presentation price
  <dbl> <chr>   <chr>      <dbl>
1 1997 japan   fresha     8204.
2 1997 japan   frozenb    8169.
3 1998 japan   fresha     7703.
4 1998 japan   frozenb    6320.
5 1999 japan   fresha     8809.
6 1999 japan   frozenb    9093.
7 2000 japan   fresha     9198.
8 2000 japan   frozenb    8557.
9 2001 japan   fresha     8260.
10 2001 japan  frozenb    5983.
# i 94 more rows

```

### Post-it down

---

<sup>3</sup>Hint: Your `names_to` argument should be a character vector of with two items. `names_sep` should be inspired by our clever use of `snake_case`.

<sup>4</sup>Hint: If you have 112 rows, remember you can use `values_drop_na = T`

### ! Values in presentation

Note that the values in the `presentation` column are not ideal. They end in `a`, `b`, `c`, and `d` due to footnotes included in Excel. For now this doesn't matter because we will quickly remove them. We'll cover some text wrangling in Week 9.

## Part E: Calculate mean annual price

### Post-it up

1. Modify the pipeline that creates `tidy_tuna_prices` to get the mean price per year<sup>5</sup>

```
# A tibble: 28 x 2
  year   price
  <dbl> <dbl>
1 1997  8186.
2 1998  7011.
3 1999  8951.
4 2000  8877.
5 2001  5633.
6 2002  5342.
7 2003  5285.
8 2004  5739.
9 2005  5554.
10 2006  5177.
# i 18 more rows
```

### Post-it down

## Exercise 2: Tidying tuna catch data (again)

### Part A: Read the tuna catch data

Note: You can copy-paste and modify your code from last week, but make sure your code is organized.

### Post-it up

1. Read in the tuna catch data from last week
2. Filter it to retain bigeye tuna (BET) caught by the purse seine fleet (PS) since 2000
3. Calculate **total** catch by year. Your final data should have 24 rows and 2 columns, as below

---

<sup>5</sup>Hint: You will use `group_by()` and `summarize()`, as well as `|>`

### **Post-it down**

```
# A tibble: 24 x 2
  year catch
  <dbl> <dbl>
1 2000 95283
2 2001 60518
3 2002 57422
4 2003 53051
5 2004 65471
6 2005 67895
7 2006 83837
8 2007 63451
9 2008 75028
10 2009 76800
# i 14 more rows
```

## **Exercise 3: Combine your catch and price data**

### **Part A: Plan the join**

1. Think about what type of join you want
2. What will be on the left and what will be on the right?
3. What is the key?
4. Write down, using human language, what you want to do.

### **Post-it up**

### **Part B: Perform the join**

1. Perform the join and save the output to an object called `tuna_revenues`
2. Create a new column that contains the annual revenue in M USD. Pay attention to the units.

### **Post-it down**

```
# A tibble: 24 x 4
  year catch price revenue
  <dbl> <dbl> <dbl>   <dbl>
1 2000 95283 8877.    846.
2 2001 60518 5633.    341.
3 2002 57422 5342.    307.
4 2003 53051 5285.    280.
5 2004 65471 5739.    376.
6 2005 67895 5554.    377.
7 2006 83837 5177.    434.
8 2007 63451 5054.    321.
```

```

9 2008 75028 5636.    423.
10 2009 76800 6175.    474.
# i 14 more rows

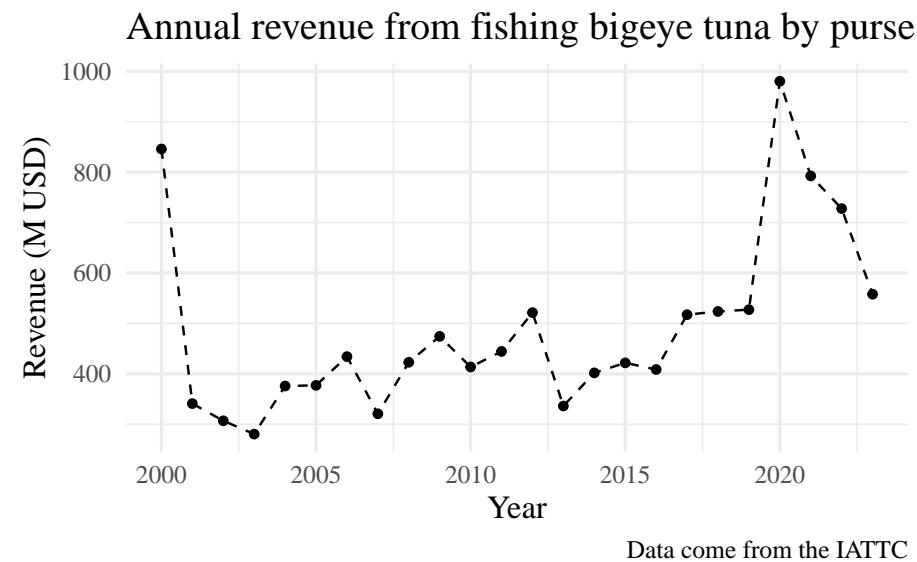
```

### Part C: Answer the questions again

1. How much TOTAL revenue since 2000?
2. How much mean ANNUAL revenue since 2000?
3. Make a figure
4. How do these plot and numbers compare to what we found [last week](#)?

[1] 11752.32

[1] 489.68



### Extra exercises for you to practice

The following four exercises use data that are built directly in R. You will need to copy and paste the provided code in your console (or R script) to make sure the objects appear in your environment.

#### Exercise 1: Pivot Longer Practice

**Scenario:** You are a TA and have been given grade data where each row represents a student and columns represent their scores on assignments 1-4. You

need to calculate the mean grade for each student.<sup>6</sup>

```
# Create the dataset
student_scores <- tribble(
  ~student_id, ~assignment_1, ~assignment_2, ~assignment_3, ~assignment_4,
  "S001", 85, 92, 78, 88,
  "S002", 91, 89, 95, 82,
  "S003", 76, 84, 91, 79,
  "S004", 88, 93, 87, 94
)

student_scores
```

# A tibble: 4 x 5  
student\_id assignment\_1 assignment\_2 assignment\_3 assignment\_4  
<chr> <dbl> <dbl> <dbl> <dbl>  
1 S001 85 92 78 88  
2 S002 91 89 95 82  
3 S003 76 84 91 79  
4 S004 88 93 87 94

### How I did it:

```
student_scores_long <- student_scores |>
  pivot_longer(cols = starts_with("assignment"),
               names_to = "assignment",
               values_to = "score") # All columns from assignment_1 to assignment_4  
# Create new column called "assignment" with the column names  
# Create new column called "score" with the values

student_means <- student_scores_long |>
  group_by(student_id) |> # Group rows by student_id
  summarize(mean_grade = mean(score)) # Calculate mean score for each student

student_means # View the result
```

# A tibble: 4 x 2  
student\_id mean\_grade  
<chr> <dbl>  
1 S001 85.8  
2 S002 89.2  
3 S003 82.5  
4 S004 90.5

---

<sup>6</sup>Hint: Use `pivot_longer()` to transform this data so that each row represents one student-assignment-score combination, then use `group_by()` and `summarize()` to calculate the mean grade for each student.

## Exercise 2: Pivot Wider Practice

**Scenario:** You have hurricane exposure data from different Florida counties. You are asked to build a figure showing the relationship between pressure and wind speed. Modify the data as needed and build a figure.<sup>7</sup>

```
# Create the dataset
hurricane_data <- tribble(
  ~county, ~metric, ~measurement,
  "Miami-Dade", "pressure", 950,
  "Miami-Dade", "precipitation", 12.5,
  "Miami-Dade", "wind_speed", 85,
  "Broward", "pressure", 955,
  "Broward", "precipitation", 8.3,
  "Broward", "wind_speed", 72,
  "Palm Beach", "pressure", 960,
  "Palm Beach", "precipitation", 6.1,
  "Palm Beach", "wind_speed", 68,
  "Monroe", "pressure", 945,
  "Monroe", "precipitation", 15.2,
  "Monroe", "wind_speed", 95
)
```

```
hurricane_data
```

```
# A tibble: 12 x 3
  county     metric      measurement
  <chr>     <chr>        <dbl>
1 Miami-Dade pressure      950
2 Miami-Dade precipitation 12.5
3 Miami-Dade wind_speed    85
4 Broward   pressure      955
5 Broward   precipitation  8.3
6 Broward   wind_speed    72
7 Palm Beach pressure      960
8 Palm Beach precipitation 6.1
9 Palm Beach wind_speed    68
10 Monroe    pressure      945
11 Monroe    precipitation 15.2
12 Monroe    wind_speed    95
```

**How I did it:**

```
hurricane_wide <- hurricane_data |>
  pivot_wider(names_from = metric,      # Use values in "metric" column as new column names
```

---

<sup>7</sup>Hint: Use `pivot_wider()` to transform this data so that each row represents a county and each metric becomes its own column.

```

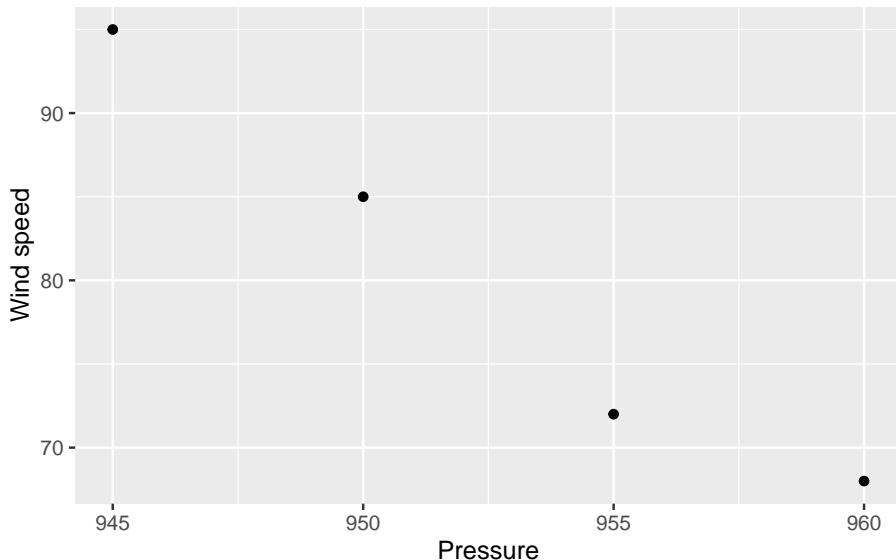
values_from = measurement) # Use values in "measurement" column to fill new columns

 # Now each county is a row with separate columns for pressure, precipitation, and wind speed

# A tibble: 4 x 4
  county      pressure precipitation wind_speed
  <chr>        <dbl>          <dbl>        <dbl>
1 Miami-Dade     950         12.5         85
2 Broward        955          8.3          72
3 Palm Beach     960          6.1          68
4 Monroe         945         15.2         95

ggplot(hurricane_wide,
       aes(x = pressure, y = wind_speed)) +
  geom_point() +
  labs(x = "Pressure",
       y = "Wind speed")

```



### Exercise 3: Joining Data Practice

**Scenario:** You have two datasets - one with student information and another with their enrollment details. You need to be able to identify the names of students taking stats courses.

```

# Create the datasets
students <- tribble(
  ~student_id, ~name, ~age,

```

```

"S001", "Alice Johnson", 20,
"S002", "Bob Smith", 22,
"S003", "Carol Davis", 19,
"S004", "David Wilson", 21,
"S005", "Eva Brown", 23
)

enrollments <- tribble(
  ~stdt_identifier, ~course, ~credits,
  "S001", "Statistics", 3,
  "S001", "Biology", 4,
  "S002", "Statistics", 3,
  "S003", "Chemistry", 4,
  "S004", "Statistics", 3,
  "S004", "Physics", 4,
  "S006", "Math", 3
)

students
```

# A tibble: 5 x 3

	student_id	name	age
1	S001	Alice Johnson	20
2	S002	Bob Smith	22
3	S003	Carol Davis	19
4	S004	David Wilson	21
5	S005	Eva Brown	23

```
enrollments
```

# A tibble: 7 x 3

	stdt_identifier	course	credits
1	S001	Statistics	3
2	S001	Biology	4
3	S002	Statistics	3
4	S003	Chemistry	4
5	S004	Statistics	3
6	S004	Physics	4
7	S006	Math	3

### How I did it:

```
combined_data <- students |>
  left_join(enrollments, by = join_by(student_id == stdt_identifier)) # Keep all students,
```

```

combined_data                                # View the combined data

# A tibble: 7 x 5
  student_id name      age course   credits
  <chr>      <chr>     <dbl> <chr>     <dbl>
1 S001       Alice Johnson  20 Statistics 3
2 S001       Alice Johnson  20 Biology    4
3 S002       Bob Smith     22 Statistics 3
4 S003       Carol Davis   19 Chemistry 4
5 S004       David Wilson  21 Statistics 3
6 S004       David Wilson  21 Physics   4
7 S005       Eva Brown    23 <NA>      NA

statistics_students <- combined_data |>
  filter(course == "Statistics") |> # Keep only rows where course equals "Statistics"
  select(name, course)               # Keep only the name and course columns

statistics_students                         # View the result

# A tibble: 3 x 2
  name      course
  <chr>     <chr>
1 Alice Johnson Statistics
2 Bob Smith    Statistics
3 David Wilson Statistics

```

### Exercise 4: Sharks!

**Scenario:** You have swimming data from beachgoers and bull shark detection data from acoustic telemetry during fourth of July. The swimming data tell you when someone entered and left the water. The shark detection data tells you which sharks were detected within the acoustic array in front of the beach, and the time of detection. Who was in the water while a shark was nearby?<sup>8</sup>

```

# Create the datasets
swimming_data <- tribble(
  ~name, ~swim_start, ~swim_end,
  "Alice", "2024-07-04 10:30:00", "2024-07-04 11:15:00",
  "Bob", "2024-07-04 10:45:00", "2024-07-04 11:30:00",
  "Carol", "2024-07-04 11:00:00", "2024-07-04 11:45:00",
  "David", "2024-07-04 11:20:00", "2024-07-04 12:00:00",
  "Eva", "2024-07-04 12:10:00", "2024-07-04 12:45:00"
) |>
  mutate(swim_start = as_datetime(swim_start),
         swim_end = as_datetime(swim_end))

```

---

<sup>8</sup>Hint: Look at the documentation for `join_by()`. What does it say about “Overlap helpers”? You’ll want to use the `between()` function.

```

shark_detections <- tribble(
  ~shark_id, ~detection_time,
  "SH001", "2024-07-04 09:40:00",
  "SH002", "2024-07-04 11:25:00",
  "SH003", "2024-07-04 11:35:00"
) |>
  mutate(detection_time = as_datetime(detection_time))

swimming_data

# A tibble: 5 x 3
  name   swim_start      swim_end
  <chr> <dttm>        <dttm>
1 Alice 2024-07-04 10:30:00 2024-07-04 11:15:00
2 Bob   2024-07-04 10:45:00 2024-07-04 11:30:00
3 Carol 2024-07-04 11:00:00 2024-07-04 11:45:00
4 David 2024-07-04 11:20:00 2024-07-04 12:00:00
5 Eva   2024-07-04 12:10:00 2024-07-04 12:45:00

shark_detections

```

### How I did it:

```

swimmer_shark_overlap <- shark_detections |>
  inner_join(swimming_data,           # Note that I am using an inner join. Play with inner
             by = join_by(between(detection_time, swim_start, swim_end))) # Find swimmers in

swimmer_shark_overlap               # View all the overlap data

# A tibble: 5 x 5
  shark_id detection_time     name   swim_start      swim_end
  <chr>    <dttm>          <chr> <dttm>        <dttm>
1 SH002    2024-07-04 11:25:00 Bob   2024-07-04 10:45:00 2024-07-04 11:30:00
2 SH002    2024-07-04 11:25:00 Carol 2024-07-04 11:00:00 2024-07-04 11:45:00
3 SH002    2024-07-04 11:25:00 David 2024-07-04 11:20:00 2024-07-04 12:00:00
4 SH003    2024-07-04 11:35:00 Carol 2024-07-04 11:00:00 2024-07-04 11:45:00
5 SH003    2024-07-04 11:35:00 David 2024-07-04 11:20:00 2024-07-04 12:00:00

at_risk_swimmers <- swimmer_shark_overlap |>
  group_by(name) |> # Keep only the columns we want to see

```

```
    summarize(n_sharks_near = n_distinct(shark_id))

at_risk_swimmers # These swimmers were in the water when sharks were detected
```

```
# A tibble: 3 x 2
  name   n_sharks_near
  <chr>      <int>
1 Bob            1
2 Carol          2
3 David          2
```